

ECE 1513: Introduction to Machine Learning

Lecture 10: Convolutional Neural Networks and Sequence Data

Ali Bereyhi

`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering
University of Toronto

Winter 2025


Recap: *Building and Training NN*

TrainingLoop():


- 1: Build NN $y = f_w(x)$ with some **hyperparameters** and **initial weights**
- 2: Split training set to mini-batches
- 3: Specify the loss function \mathcal{L}
- 4: **for** $epochs = 1, \dots, E$ **do**
- 5: Keep applying mini-batch SGD
- 6: **end for**
- 7: Return final weights w^* , average training loss, and accuracy on training set

Recap: Generalization

- ? *How can we measure generalization?*
- ! *We try samples that we did not use for training*



Recap: Underfitting and Overfitting



- Left \rightsquigarrow Underfitting
- Middle \rightsquigarrow Learning fairly
- Right \rightsquigarrow Overfitting

Recap: *Bias and Variance*

Generalization Error Components


Generalization error is proportional to the model bias and variance

We can only minimize the bias and variance of our model output



- *If we set the model to give us zero bias \rightsquigarrow unbiased estimator*
 - ↳ *But this may lead to higher variance*
- *There is always a minimum error that we cannot beat \rightsquigarrow Bayes error*

Recap: Convolution

In convolution, we slide a filter over the input sample



Recap: Convolution as Sparse Linear Transform



Today's Agenda: CNNs and Sequence Data

Today, we use the convolution to build more efficient FNNs called

Convolutional Neural Networks

In this way, we learn

- *Convolutional layers*
- *Pooling layers*
- *Architecture of deep CNNs*

We then start our final journey which explores briefly


Advanced Topics in Machine Learning

In today's episode, we talk about


- *Sequence Data and Seq2Seq Models*

Building Neural Layers with Convolution

We can activate the convolution output to make a partially-connected layer




Convolutional Layer: *Single-Channel Output*



Convolutional Layer: *Multi-Channel Output*

To have multiple output channels, we can use multiple filters



Pooling: Max-Pooling

Pooling

Pooling is a convolution-like operation that computes a fix function in each slide


Example: Max-Pooling

$$\mathbf{X} = \begin{bmatrix} X_{1,1} & X_{1,2} & \dots & \dots & X_{1,M} \\ X_{2,1} & X_{2,2} & \dots & \dots & X_{2,M} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ X_{N,1} & X_{N,2} & \dots & \dots & X_{N,M} \end{bmatrix} \rightarrow Z = \max \{\text{window}\}$$

we pool the *maximum*

$$\mathbf{Z} = \begin{bmatrix} Z_{1,1} & Z_{1,2} & \dots & \dots \\ \dots & \dots & \dots & \dots \\ \dots & \dots & \dots & \dots \end{bmatrix}$$

Max Pooling: Numerical Example



3.2	3.2	2.1	1.5
3.2	3.8	3.8	4.3
2.3	3.8	3.8	4.3

Mean-Pooling

Mean-pooling is another approach in which we compute the average


$$\mathbf{X} = \begin{bmatrix} X_{1,1} & X_{1,2} & \dots & \dots & X_{1,M} \\ X_{2,1} & X_{2,2} & \dots & \dots & X_{2,M} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ X_{N,1} & X_{N,2} & \dots & \dots & X_{N,M} \end{bmatrix}$$

mean { }

In each window, we pool the *average*

$$\mathbf{Z} = \begin{bmatrix} Z_{1,1} & Z_{1,2} & \dots & \dots \\ \dots & \dots & \dots & \dots \\ \dots & \dots & \dots & \dots \end{bmatrix}$$

Mean-Pooling: Numerical Example




2.05	1.9	1.225	0.7
1.85	2.275	1.85	1.45
1.525	1.9	1.5	1.775


Convolution with Stride

We can perform all operations with *stride*, e.g., *stride 2*


input



filter




output



Convolutional Unit: Complex Convolutional Layer


A convolutional unit is made as



- *The convolution performs less complex linear operation*
- *Pooling make the output smooth, i.e., less varying*

CNN: General Architecture


A CNN consists of multiple convolutional units and a fully-connected NN




We can repeat the convolutional unit multiple times

Example: AlexNet

Let's look at the winner of the ImageNet challenge in 2012




Example: Custom CNN for MNIST Classification



In this network, we do the following

- We apply convolution with 10 filters
 - ↳ We apply pooling with stride 2
- We apply convolution with 20 filters
 - ↳ We apply pooling with stride 2
- We flatten and pass through one fully-connected layer
 - ↳ We compute Softmax for classification

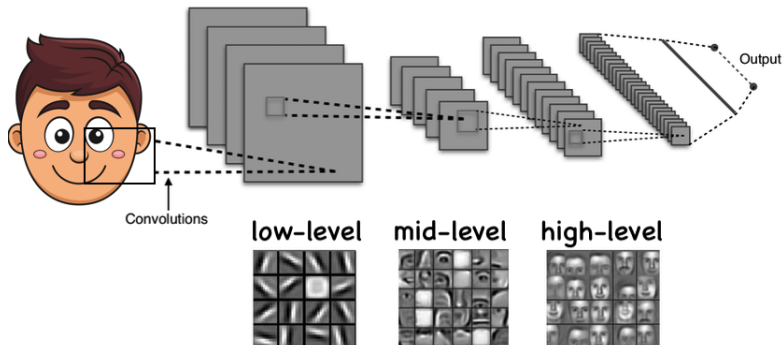
Backpropagation



It is easy to see that backward pass is dual to the forward pass


- *We can compute sample gradients with backpropagation*
- *Using mini-batch SGD, we can efficiently train CNNs*

Gradual Feature Extraction



As we go deeper, CNN extracts higher levels of features

Vanishing or Exploding Gradient




As we go very deep in FNNs, we could experience

- Decrease in gradient values through depth
 - ↳ This results in **vanishing** gradient
- Increase in gradient values through depth
 - ↳ This results in **exploding** gradient

ResNet: Skip Connection

ResNet uses *skip connections* to overcome this issue



Further Read

- Goodfellow
 - ↳ Chapter 9

CNNs

CNNs are further discussed in details in

- *ECE1508: Applied Deep Learning*
 - ↳ *Given in both Fall and Winter Semesters*

Sequence Data: *Many Applications*



"This is the first lecture on ..."



Classic/Jazz/Pop/Hip-hop

"This product is useless!"



A cute puppy in snow



Sport/Drama/Documentary

Sequence Learning Problem

Basic FNNs cannot be used in practice: *we need a huge input and/or output*



Sequence Learning Model

Models that get sequence inputs and return sequence outputs




There are various approaches in the literature

- *Recurrent Neural Networks (RNNs)*
- *Encoder-Decoder Architecture (Seq2Seq Models)*
- *Transformers*

Types of Sequence Problems



many-to-many




Recurrent Neural Networks

We can feed a sequence to a neural network, if we include **recurrence**


Recurrence

NN takes a state as input and returns the next state as output



Example: *Turning Shallow FNN to Basic RNN*

Say we have a shallow fully-connected FNN




- *It computes hidden feature from the input x*

$$\mathbf{h} = f(\mathbf{W}_1 \mathbf{x})$$

- *It computes the output from the hidden features*

$$\mathbf{y} = f(\mathbf{W}_2 \mathbf{h})$$

Example: *Turning Shallow FNN to Basic RNN*



The model in this case gives us

$$\mathbf{y} \propto P(v|\mathbf{x})$$

and when we train it, we maximize its likelihood

Example: *Turning Shallow FNN to Basic RNN*

We can make an RNN by using the previous feature in each time: at time t


- compute new features from the input $x[t]$ and previous features $\mathbf{h}[t - 1]$*

$$\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_0 \mathbf{h}[t - 1])$$

- computes the output from the new features*

$$\mathbf{y}[t] = f(\mathbf{W}_2 \mathbf{h}[t])$$

Example: *Turning Shallow FNN to Basic RNN*




This recursive model determines for us

$$y[t] \propto P(v[t] | h[t-1], x[t]) \equiv P(v[t] | x[1], \dots, x[t])$$

and we should maximize the likelihood over time


↳ *information about $x[1], \dots, x[t-1]$ is somehow encoded in $h[t-1]$*

Example: *Turning Shallow FNN to Basic RNN*



Generic RNN


We can do the same thing with any FNN, including CNNs



$h[t]$ is the output of a ny number of hidden layers

Training RNNs

We can apply the same risk minimization; this time over time



To train RNNs, we need to backpropagate through time

Main Challenge: *Vanishing Gradient Through Time*

With long sequences

backpropagation through time \approx *backpropagation through* **deep** NNs


- *Exploding Gradients*
- *Vanishing Gradients*
 - ↳ *Gradients become small over time*
 - ↳ *Weights of RNN are updated only with most recent time instances*
 - ↳ *RNN has a limited memory!*

Classical Remedies

There are various approaches to handle these issues


- *Clipping gradients when exploding*
 - ↳ *Usually used to deal with exploding gradient*
- *Truncated backpropagation through time*
 - ↳ *Updated multiple times in the middle to carry memory forward*
- *Using gated units*
 - ↳ *Use the so-called gate to control better the memory*

Basic RNN as a Unit




Gated Recurrent Unit: GRU

This is what's going on in a GRU cell




Gated Recurrent Unit: GRU

Compute *update gate* $\Gamma_u[t] = \sigma(\mathbf{W}_{u,\text{in}}\mathbf{x}[t] + \mathbf{W}_{u,\text{m}}\mathbf{h}[t-1])$




Gated Recurrent Unit: GRU

Compute *reset gate* $\Gamma_r[t] = \sigma(\mathbf{W}_{r,\text{in}}\mathbf{x}[t] + \mathbf{W}_{r,\text{m}}\mathbf{h}[t-1])$




Gated Recurrent Unit: GRU

Compute *actual memory* $\tilde{\mathbf{h}}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{\Gamma}_r[t] \odot \mathbf{h}[t - 1])$




Gated Recurrent Unit: GRU

Update *hidden state* as $\mathbf{h}[t] = (1 - \mathbf{\Gamma}_u[t]) \odot \mathbf{h}[t-1] + \mathbf{\Gamma}_u[t] \odot \tilde{\mathbf{h}}[t]$




Long Short-Term Memory: *LSTM*

This is how inside an *LSTM* unit looks like




Long Short-Term Memory: *LSTM*

Actual cell state $\tilde{\mathbf{c}}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$




Long Short-Term Memory: *LSTM*

We use *forget gate* and *update gate* to update cell state



Long Short-Term Memory: *LSTM*

We use *output gate* to control fellow of memory to the *hidden state*



Gating is Helpful

- *Using gating we can control vanishing and exploding gradient*
- *LSTM used to be a robust sequence-based model*
 - ↳ *We can still use it for simple tasks like price prediction*
 - ↳ *It was one of the first models used for text generation*
- *At the end of the day, RNNs will still carry limited memory*
 - ↳ *This is why **Attention Mechanism** is developed*
 - ↳ *We can do whole sequence processing based on **Attention***
 - ↳ *This is the idea of **Transformers***

Further Read

- Goodfellow
 - ↳ Chapter 10

RNNs

RNNs and Transformers are discussed in

- *ECE1508: Applied Deep Learning*
 - ↳ *Given in both Fall and Winter Semesters*
- *ECE1786: Creative Applications of NLP*
 - ↳ *Given in Fall Semesters*